

2020_1125_Comparison_with_Linear_Models

November 25, 2020

1 Predicting Concrete Compressive Strength - Comparison with Linear Models

In this code notebook, we will analyze the statistics pertaining the various models presented in this project. In the Exploratory Data Analysis notebook, we explored the various relationships that each constituent of concrete has on the cured compressive strength. The materials that held the strongest relationships, regardless of curing time, were cement, cementitious ratio, superplasticizer ratio, and fly ash ratio. We will examine each of the linear ratios independent of age, as well as at the industry-standard 28 day cure time mark.

1.1 Dataset Citation

This dataset was retrieved from the UC Irvine Machine Learning Repository from the following URL: <https://archive.ics.uci.edu/ml/datasets/Concrete+Compressive+Strength>.

The dataset was donated to the UCI Repository by Prof. I-Cheng Yeh of Chung-Huah University, who retains copyright for the following published paper: I-Cheng Yeh, “Modeling of strength of high performance concrete using artificial neural networks,” Cement and Concrete Research, Vol. 28, No. 12, pp. 1797-1808 (1998). Additional papers citing this dataset are listed at the reference link above.

1.2 Import the Relevant Libraries

```
[1]: # Data Manipulation
import numpy as np
import pandas as pd

# Data Visualization
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set()

# Data Preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

# Linear Regresssion Model
```

```

from sklearn.linear_model import LinearRegression

# Model Evaluation
from sklearn.metrics import \
    mean_squared_error, mean_absolute_error, explained_variance_score

```

1.3 Import & Check the Data

```

[2]: df1 = pd.read_csv('2020_1124_Modeling_Data.csv')
     df2 = pd.read_csv('2020_1123_Concrete_Data_Loaded_Transformed.csv')

     original_data = df1.copy()
     transformed_data = df2.copy()

```

```

[3]: # The original data contains kg/m^3 values
     original_data.head()

```

```

[3]:
    Cement  Blast_Furnace_Slag  Fly_Ash  Water  Superplasticizer  \
0    540.0                0.0      0.0  162.0                2.5
1    540.0                0.0      0.0  162.0                2.5
2    332.5               142.5      0.0  228.0                0.0
3    332.5               142.5      0.0  228.0                0.0
4    198.6               132.4      0.0  192.0                0.0

    Coarse_Aggregate  Fine_Aggregate  Age  Compressive_Strength
0             1040.0           676.0   28              79.99
1             1055.0           676.0   28              61.89
2              932.0           594.0  270              40.27
3              932.0           594.0  365              41.05
4              978.4           825.5  360              44.30

```

```

[4]: # Original data
     original_data.describe()

```

```

[4]:
    count  Cement  Blast_Furnace_Slag  Fly_Ash  Water  \
count  1030.000000      1030.000000  1030.000000  1030.000000
mean    281.167864      73.895825    54.188350   181.567282
std     104.506364      86.279342    63.997004    21.354219
min     102.000000       0.000000     0.000000   121.800000
25%     192.375000       0.000000     0.000000   164.900000
50%     272.900000      22.000000     0.000000   185.000000
75%     350.000000     142.950000   118.300000   192.000000
max      540.000000     359.400000    200.100000   247.000000

    Superplasticizer  Coarse_Aggregate  Fine_Aggregate  Age  \
count    1030.000000      1030.000000    1030.000000  1030.000000
mean         6.204660      972.918932      773.580485    45.662136

```

std	5.973841	77.753954	80.175980	63.169912
min	0.000000	801.000000	594.000000	1.000000
25%	0.000000	932.000000	730.950000	7.000000
50%	6.400000	968.000000	779.500000	28.000000
75%	10.200000	1029.400000	824.000000	56.000000
max	32.200000	1145.000000	992.600000	365.000000

	Compressive_Strength
count	1030.000000
mean	35.817961
std	16.705742
min	2.330000
25%	23.710000
50%	34.445000
75%	46.135000
max	82.600000

```
[5]: # The transformed data contains ratios to total mass of the concrete mix
transformed_data.head()
```

```
[5]:   Cementitious_Ratio  Slag_Ratio  Fly_Ash_Ratio  Water_to_Cementitious_Ratio  \
0          0.205086      0.000000           0.0          0.400000
1          0.167391      0.000000           0.0          0.483117
2          0.058291      0.087436           0.0          1.375358
3          0.145726      0.000000           0.0          0.550143
4          0.085350      0.056900           0.0          0.966767
```

	Superplasticizer_Ratio	Coarse_Aggregate_Ratio	Sand_Ratio	Age	\
0	0.0	0.461444	0.251436	1	
1	0.0	0.420000	0.331739	1	
2	0.0	0.437179	0.336924	3	
3	0.0	0.437179	0.336924	3	
4	0.0	0.420474	0.354764	3	

	Compressive_Strength
0	12.638095
1	6.267337
2	8.063422
3	15.049193
4	9.131420

```
[6]: # Transformed data
transformed_data.describe()
```

```
[6]:   Cementitious_Ratio  Slag_Ratio  Fly_Ash_Ratio  \
count          1030.000000    1030.000000    1030.000000
mean             0.142726      0.031643      0.155263
```

std	0.040513	0.036961	0.187884
min	0.044815	0.000000	0.000000
25%	0.124002	0.000000	0.000000
50%	0.143272	0.009455	0.000000
75%	0.162794	0.061972	0.319960
max	0.259517	0.150339	0.588415

	Water_to_Cementitious_Ratio	Superplasticizer_Ratio \
count	1030.000000	1030.000000
mean	0.611796	0.002620
std	0.278319	0.002494
min	0.265918	0.000000
25%	0.447540	0.000000
50%	0.547837	0.002727
75%	0.666639	0.004338
max	1.882353	0.013149

	Coarse_Aggregate_Ratio	Sand_Ratio	Age	Compressive_Strength
count	1030.000000	1030.000000	1030.000000	1030.000000
mean	0.415166	0.330117	45.662136	35.817836
std	0.031021	0.033245	63.169912	16.705679
min	0.345890	0.247971	1.000000	2.331808
25%	0.392294	0.311208	7.000000	23.707115
50%	0.420464	0.330543	28.000000	34.442774
75%	0.437623	0.354096	56.000000	46.136287
max	0.479846	0.414147	365.000000	82.599225


1.4 Cement Modeling - Including All Cure Times

We understand that the ratio of cement to compressive strength is linear. We will model this relationship in Python and evaluate its performance compared to our ANN model.

1.4.1 Visualization

```
[7]: # We will visualize the linear relationship between quantity of cement and ↵
      ↪ compressive strength
      cement = original_data['Cement']
      strength = original_data['Compressive_Strength']
      plt.scatter(cement, strength)
```

```
[7]: <matplotlib.collections.PathCollection at 0x7ff55c9991c0>
```



1.4.2 Train the Linear Model

```
[8]: # Reshape the data so it complies with the linear model requirements  
X = np.array(cement).reshape(1030,1)  
y = np.array(strength).reshape(1030,1)
```

```
[9]: # Perform a train-test split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.  
↪2,random_state=42)  
  
# Train the linear model  
lm = LinearRegression()  
lm.fit(X_train,y_train)
```

```
[9]: LinearRegression()
```

1.4.3 Test the Linear Model

```
[10]: y_pred = lm.predict(X_test)
```

1.4.4 Linear Equation

```
[11]: # print the intercept  
print(lm.intercept_)
```

```
[13.78517188]
```


```
[12]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])  
coeff
```

```
[12]:      Coefficient  
0      0.077896
```

1.4.5 Model Evaluation

```
[13]: # Plot the linear model predictions as a line superimposed on a scatter plot of the testing data  
plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

```
[13]: [<matplotlib.lines.Line2D at 0x7ff55cd19760>]
```



```
[14]: # Evaluation Metrics  
MAE_cement = mean_absolute_error(y_test, y_pred)  
MSE_cement = mean_squared_error(y_test, y_pred)  
RMSE_cement = np.sqrt(mean_squared_error(y_test, y_pred))
```

```

cement_stats = [MAE_cement,MSE_cement,RMSE_cement] # storing for model
↳ comparison at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR CEMENT VS. COMPRESSIVE STRENGTH")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_cement}\nMean Squared Error:
↳\t\t\t{MSE_cement}\nRoot Mean Squared Error (RMSE):\t\t{RMSE_cement}")
print('-----\n\n')

```

EVALUATION METRICS, LINEAR MODEL FOR CEMENT VS. COMPRESSIVE STRENGTH

```

-----
Mean Absolute Error (MAE):                11.55561279863471
Mean Squared Error:                      192.78479855432548
Root Mean Squared Error (RMSE):          13.884696559677689
-----

```

1.5 Cement Modeling - 28 Day Cure Time

We will model the cement vs compressive strength relationship for a constant cure time (28 days).

1.5.1 Visualization

```


[15]: # We will visualize the linear relationship between quantity of cement and
↳ compressive strength at 28 days
cement = original_data[original_data['Age']==28]['Cement']
strength = original_data[original_data['Age']==28]['Compressive_Strength']
plt.scatter(cement,strength)

```

```

[15]: <matplotlib.collections.PathCollection at 0x7ff55cd70d90>

```



1.5.2 Train the Linear Model

```
[16]: # Reshape the data so it complies with the linear model requirements  
X = np.array(cement).reshape(425,1)  
y = np.array(strength).reshape(425,1)
```

```
[17]: # Perform a train-test split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.  
↪2,random_state=42)  
  
# Train the linear model  
lm = LinearRegression()  
lm.fit(X_train,y_train)
```

```
[17]: LinearRegression()
```

1.5.3 Test the Linear Model

```
[18]: y_pred = lm.predict(X_test)
```


1.5.4 Linear Equation

```
[19]: # print the intercept  
print(lm.intercept_)
```

[13.07410297]


```
[20]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])  
coeff
```

```
[20]:      Coefficient  
0      0.088248
```

1.5.5 Model Evaluation

```
[21]: # Plot the linear model predictions as a line superimposed on a scatter plot of u  
      ↪ the testing data  
plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

```
[21]: [<matplotlib.lines.Line2D at 0x7ff55dbb01f0>]
```



```
[22]: # Evaluation Metrics  
MAE_cement_28 = mean_absolute_error(y_test, y_pred)  
MSE_cement_28 = mean_squared_error(y_test, y_pred)  
RMSE_cement_28 = np.sqrt(mean_squared_error(y_test, y_pred))
```

```

cement_28_stats = [MAE_cement_28,MSE_cement_28,RMSE_cement_28] # storing for
↳model comparison at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR CEMENT VS. COMPRESSIVE STRENGTH")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_cement_28}\nMean Squared Error:
↳\t\t\t{MSE_cement_28}\nRoot Mean Squared Error (RMSE):\t\t{RMSE_cement_28}")
print('-----\n\n')

```

EVALUATION METRICS, LINEAR MODEL FOR CEMENT VS. COMPRESSIVE STRENGTH

```

-----
Mean Absolute Error (MAE):          9.134081632197555
Mean Squared Error:                140.11750347700305
Root Mean Squared Error (RMSE):    11.837123952928899
-----

```

1.6 Cementitious Ratio Modeling - Including All Cure Times

We know that the ratio of cementitious materials to the total mass is (cement + fly ash)/(total mass) to compressive strength is linear. We will model this relationship in Python and evaluate its performance.


1.6.1 Visualization

```

[23]: # We will visualize the linear relationship between quantity of cementitious
↳materials and compressive strength
cementitious = transformed_data['Cementitious_Ratio']
strength = transformed_data['Compressive_Strength']
plt.scatter(cementitious,strength)

```

[23]: <matplotlib.collections.PathCollection at 0x7ff55dbe5c10>



1.6.2 Train the Linear Model

```
[24]: # Reshape the data so it complies with the linear model requirements
X = np.array(cementitious).reshape(1030,1)
y = np.array(strength).reshape(1030,1)
```

```
[25]: # Perform a train-test split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
↪2,random_state=42)

# Train the linear model
lm = LinearRegression()
lm.fit(X_train,y_train)
```

```
[25]: LinearRegression()
```

1.6.3 Test the Linear Model

```
[26]: y_pred = lm.predict(X_test)
```

1.6.4 Linear Equation

```
[27]: # print the intercept  
print(lm.intercept_)
```

```
[10.00280955]
```


```
[28]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])  
coeff
```

```
[28]:      Coefficient  
0      180.19044
```

1.6.5 Model Evaluation

```
[29]: # Plot the linear model predictions as a line superimposed on a scatter plot of the testing data  
plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

```
[29]: [<matplotlib.lines.Line2D at 0x7ff55dd229a0>]
```



```
[30]: # Evaluation Metrics  
MAE_cementitious = mean_absolute_error(y_test, y_pred)  
MSE_cementitious = mean_squared_error(y_test, y_pred)  
RMSE_cementitious = np.sqrt(mean_squared_error(y_test, y_pred))
```

```

cementitious_stats = [MAE_cementitious,MSE_cementitious,RMSE_cementitious] #
    ↳storing for model comparison at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR CEMENTITIOUS RATIO VS. COMPRESSIVE_
    ↳STRENGTH")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_cementitious}\nMean Squared Error:
    ↳\t\t\t{MSE_cementitious}\nRoot Mean Squared Error (RMSE):
    ↳\t\t\t{RMSE_cementitious}")
print('-----\n\n')

```

EVALUATION METRICS, LINEAR MODEL FOR CEMENTITIOUS RATIO VS. COMPRESSIVE STRENGTH

```

-----
Mean Absolute Error (MAE):                12.834672509172698
Mean Squared Error:                      253.05814774985234
Root Mean Squared Error (RMSE):          15.907801474429217
-----

```

1.7 Cementitious Ratio Modeling - 28 Day Cure Time


1.7.1 Visualization

```

[31]: # We will visualize the linear relationship between quantity of cementitious_
    ↳materials and compressive strength at 28 days
cementitious = transformed_data[original_data['Age']==28]['Cementitious_Ratio']
strength = transformed_data[original_data['Age']==28]['Compressive_Strength']
plt.scatter(cementitious,strength)

```

[31]: <matplotlib.collections.PathCollection at 0x7ff55de47fa0>



1.7.2 Train the Linear Model

```
[32]: # Reshape the data so it complies with the linear model requirements  
X = np.array(cementitious).reshape(425,1)  
y = np.array(strength).reshape(425,1)
```

```
[33]: # Perform a train-test split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.  
↪2,random_state=42)  
  
# Train the linear model  
lm = LinearRegression()  
lm.fit(X_train,y_train)
```

```
[33]: LinearRegression()
```

1.7.3 Test the Linear Model

```
[34]: y_pred = lm.predict(X_test)
```

1.7.4 Linear Equation

```
[35]: # print the intercept  
print(lm.intercept_)
```

```
[13.59333113]
```


```
[36]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])  
coeff
```

```
[36]:      Coefficient  
0      191.571248
```

1.7.5 Model Evaluation

```
[37]: # Plot the linear model predictions as a line superimposed on a scatter plot of u  
      ↪ the testing data  
plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

```
[37]: [<matplotlib.lines.Line2D at 0x7ff55deaa880>]
```



```
[38]: # Evaluation Metrics  
MAE_cementitious_28 = mean_absolute_error(y_test, y_pred)  
MSE_cementitious_28 = mean_squared_error(y_test, y_pred)  
RMSE_cementitious_28 = np.sqrt(mean_squared_error(y_test, y_pred))
```

```

cementitious_28_stats =
    ↳ [MAE_cementitious_28, MSE_cementitious_28, RMSE_cementitious_28] # storing for
    ↳ model comparison at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR CEMENTITIOUS RATIO VS. COMPRESSIVE_
    ↳ STRENGTH AT 28 DAYS")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_cementitious_28}\nMean Squared Error:
    ↳ \t\t\t{MSE_cementitious_28}\nRoot Mean Squared Error (RMSE):
    ↳ \t\t\t{RMSE_cementitious_28}")
print('-----\n\n')

```

EVALUATION METRICS, LINEAR MODEL FOR CEMENTITIOUS RATIO VS. COMPRESSIVE STRENGTH
AT 28 DAYS

```

-----
Mean Absolute Error (MAE):                11.519580245958837
Mean Squared Error:                      197.95720060481753
Root Mean Squared Error (RMSE):          14.069726386991949
-----

```

1.8 Fly Ash Ratio Modeling - Including All Cure Times

The fly ash ratio is interpreted as the percentage of fly ash within the cementitious materials mix, that is, $\text{Fly_Ash_Ratio} = (\text{fly ash} + \text{cement}) / (\text{total mass})$.


1.8.1 Visualization

```

[39]: # We will visualize the linear relationship between fly ash ratio and
    ↳ compressive strength
fly = transformed_data['Fly_Ash_Ratio']
strength = transformed_data['Compressive_Strength']
plt.scatter(fly, strength)

```

[39]: <matplotlib.collections.PathCollection at 0x7ff55db156a0>




1.8.2 Data Preprocessing

We see from the graph above that there are many instances where there is no fly ash in the mix design. Let us use only nonzero entries for our analysis.

```
[40]: fly = transformed_data[transformed_data['Fly_Ash_Ratio']!=0]['Fly_Ash_Ratio']
      strength = transformed_data[transformed_data['Fly_Ash_Ratio']!=
      ↪=0]['Compressive_Strength']
      plt.scatter(fly,strength)
```

```
[40]: <matplotlib.collections.PathCollection at 0x7ff5600d4a60>
```



1.8.3 Train the Linear Model

```
[41]: # Reshape the data so it complies with the linear model requirements  
X = np.array(fly).reshape(464,1)  
y = np.array(strength).reshape(464,1)
```

```
[42]: # Perform a train-test split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.  
↪2,random_state=42)  
  
# Train the linear model  
lm = LinearRegression()  
lm.fit(X_train,y_train)
```

```
[42]: LinearRegression()
```

1.8.4 Test the Linear Model

```
[43]: y_pred = lm.predict(X_test)
```

1.8.5 Linear Equation

```
[44]: # print the intercept  
print(lm.intercept_)
```

```
[53.13047155]
```


```
[45]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])  
coeff
```

```
[45]:    Coefficient  
0    -54.508785
```

1.8.6 Model Evaluation

```
[46]: # Plot the linear model predictions as a line superimposed on a scatter plot of the testing data  
plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

```
[46]: [<matplotlib.lines.Line2D at 0x7ff56012ddc0>]
```



```
[47]: # Evaluation Metrics  
MAE_fly = mean_absolute_error(y_test, y_pred)  
MSE_fly = mean_squared_error(y_test, y_pred)  
RMSE_fly = np.sqrt(mean_squared_error(y_test, y_pred))
```

```

fly_stats = [MAE_fly,MSE_fly,RMSE_fly] # storing for model comparison at the
↳end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR FLY ASH RATIO VS. COMPRESSIVE_
↳STRENGTH")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_fly}\nMean Squared Error:
↳\t\t\t{MSE_fly}\nRoot Mean Squared Error (RMSE):\t\t{RMSE_fly}")
print('-----\n\n')

```

```

EVALUATION METRICS, LINEAR MODEL FOR FLY ASH RATIO VS. COMPRESSIVE STRENGTH
-----
Mean Absolute Error (MAE):                12.121987321537818
Mean Squared Error:                      212.8943249027957
Root Mean Squared Error (RMSE):          14.590898700998363
-----

```

1.9 Fly Ash Ratio Modeling - 28 Day Cure Time

The fly ash ratio is interpreted as the percentage of fly ash within the cementitious materials mix, that is, $\text{Fly_Ash_Ratio} = (\text{fly ash} + \text{cement}) / (\text{total mass})$.

```


[48]: fly = transformed_data[((transformed_data['Fly_Ash_Ratio']!
↳=0)&(transformed_data['Age']==28))]['Fly_Ash_Ratio']
strength = transformed_data[((transformed_data['Fly_Ash_Ratio']!
↳=0)&(transformed_data['Age']==28))]['Compressive_Strength']
plt.scatter(fly,strength)

```

```

[48]: <matplotlib.collections.PathCollection at 0x7ff5601eac40>

```



1.9.1 Train the Linear Model

```
[49]: # Reshape the data so it complies with the linear model requirements
X = np.array(fly).reshape(217,1)
y = np.array(strength).reshape(217,1)
```

```
[50]: # Perform a train-test split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
↪2,random_state=42)

# Train the linear model
lm = LinearRegression()
lm.fit(X_train,y_train)
```

```
[50]: LinearRegression()
```

1.9.2 Test the Linear Model

```
[51]: y_pred = lm.predict(X_test)
```

1.9.3 Linear Equation

```
[52]: # print the intercept  
print(lm.intercept_)
```

[60.04189914]


```
[53]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])  
coeff
```

```
[53]:      Coefficient  
0      -71.730364
```

1.9.4 Model Evaluation

```
[54]: # Plot the linear model predictions as a line superimposed on a scatter plot of u  
      ↪ the testing data  
plt.scatter(X_test,y_test)  
plt.plot(X_test,y_pred,'r')
```

```
[54]: [<matplotlib.lines.Line2D at 0x7ff56036bd00>]
```



```
[55]: # Evaluation Metrics  
MAE_fly_28 = mean_absolute_error(y_test, y_pred)  
MSE_fly_28 = mean_squared_error(y_test, y_pred)  
RMSE_fly_28 = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
fly_28_stats = [MAE_fly_28,MSE_fly_28,RMSE_fly_28] # storing for model
↳ comparison at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR FLY ASH RATIO VS. COMPRESSIVE_
↳ STRENGTH AT 28 DAYS")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_fly_28}\nMean Squared Error:
↳ \t\t\t{MSE_fly_28}\nRoot Mean Squared Error (RMSE):\t\t{RMSE_fly_28}")
print('-----\n\n')
```

EVALUATION METRICS, LINEAR MODEL FOR FLY ASH RATIO VS. COMPRESSIVE STRENGTH AT 28 DAYS

```
-----
Mean Absolute Error (MAE):          7.692694771235926
Mean Squared Error:                91.3179595195656
Root Mean Squared Error (RMSE):    9.556043089038768
-----
```


1.10 Superplasticizer Ratio Modeling - Including All Cure Times

The superplasticizer ratio is the ratio of superplasticizer contained within the total mix design, by weight.

1.10.1 Visualization

```
[56]: # We will visualize the linear relationship between superplasticizer ratio and
↳ compressive strength
superplasticizer = transformed_data['Superplasticizer_Ratio']
strength = transformed_data['Compressive_Strength']
plt.scatter(superplasticizer,strength)
```

```
[56]: <matplotlib.collections.PathCollection at 0x7ff5603c2310>
```




1.10.2 Data Preprocessing

Once again, we see from the graph above that there are many instances where there is no superplasticizer in the mix design. Let us use only nonzero entries for our analysis.

```
[57]: superplasticizer = transformed_data[transformed_data['Superplasticizer_Ratio']!
      ↪=0]['Superplasticizer_Ratio']
      strength = transformed_data[transformed_data['Superplasticizer_Ratio']!
      ↪=0]['Compressive_Strength']
      plt.scatter(superplasticizer,strength)
```

```
[57]: <matplotlib.collections.PathCollection at 0x7ff560526130>
```

This is better, but we see a large spread in the data. Let's remove any outliers first, before training our model.

```
[58]: superplasticizer.describe()
```

```
[58]: count    651.000000
      mean      0.004146
      std      0.001875
      min      0.000746
      25%      0.002947
      50%      0.003998
      75%      0.004834
      max      0.013149
      Name: Superplasticizer_Ratio, dtype: float64
```

```
[59]: mean = 0.004146
      three_sigma = 3*0.001875
      upper = mean + three_sigma
      lower = mean - three_sigma

      print(f"The lower bound is:\t{lower}\n\nThe upper bound is:\t{upper}")
```

```
The lower bound is:    -0.0014789999999999994
The upper bound is:    0.009771
```

Since there are no negative ratios, we only need to remove data points where the superplasticizer ratio is greater than 0.009771.

```
[60]: superplasticizer = transformed_data[transformed_data['Superplasticizer_Ratio']!=0][transformed_data['Superplasticizer_Ratio'] < upper]['Superplasticizer_Ratio']
strength = transformed_data[transformed_data['Superplasticizer_Ratio']!=0][transformed_data['Superplasticizer_Ratio'] < upper]['Compressive_Strength']
plt.scatter(superplasticizer, strength)
```


<ipython-input-60-1ad850f1b580>:1: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
superplasticizer = transformed_data[transformed_data['Superplasticizer_Ratio']!=0][transformed_data['Superplasticizer_Ratio'] < upper]['Superplasticizer_Ratio']
```

<ipython-input-60-1ad850f1b580>:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
strength = transformed_data[transformed_data['Superplasticizer_Ratio']!=0][transformed_data['Superplasticizer_Ratio'] < upper]['Compressive_Strength']
```

[60]: <matplotlib.collections.PathCollection at 0x7ff5606067c0>



1.10.3 Train the Linear Model

```
[61]: # We will train and test our model only on the data above, that does not contain outliers
# Reshape the data so it complies with the linear model requirements
```

```
X = np.array(superplasticizer).reshape(641,1)
y = np.array(strength).reshape(641,1)
```

```
[62]: # Perform a train-test split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
↪2,random_state=42)

# Train the linear model
lm = LinearRegression()
lm.fit(X_train,y_train)
```

```
[62]: LinearRegression()
```

1.10.4 Test the Linear Model

```
[63]: y_pred = lm.predict(X_test)
```

1.10.5 Linear Equation

```
[64]: # print the intercept
print(lm.intercept_)
```

```
[30.85660295]
```


```
[65]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])
coeff
```

```
[65]: Coefficient
0    2121.483927
```

1.10.6 Model Evaluation

```
[66]: # Plot the linear model predictions as a line superimposed on a scatter plot of ↵
↪the testing data
plt.scatter(X_test,y_test)
plt.plot(X_test,y_pred,'r')
```

```
[66]: [<matplotlib.lines.Line2D at 0x7ff560663250>]
```



```
[67]: # Evaluation Metrics
MAE_super = mean_absolute_error(y_test, y_pred)
MSE_super = mean_squared_error(y_test, y_pred)
RMSE_super = np.sqrt(mean_squared_error(y_test, y_pred))

super_stats = [MAE_super, MSE_super, RMSE_super] # storing for model comparison
↳ at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR SUPERPLASTICIZER RATIO VS. COMPRESSIVE STRENGTH")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_super}\nMean Squared Error:\t\t\t\t{MSE_super}\nRoot Mean Squared Error (RMSE):\t\t\t\t{RMSE_super}")
print('-----\n\n')
```

EVALUATION METRICS, LINEAR MODEL FOR SUPERPLASTICIZER RATIO VS. COMPRESSIVE STRENGTH

```
-----
Mean Absolute Error (MAE):                12.189766983759341
Mean Squared Error:                      225.90929843844432
Root Mean Squared Error (RMSE):          15.030279386573103
-----
```


1.11 Superplasticizer Ratio Modeling - 28 Day Cure Time

The superplasticizer ratio is the ratio of superplasticizer contained within the total mix design, by weight.

1.11.1 Visualization

```
[68]: superplasticizer =  
    →transformed_data[((transformed_data['Superplasticizer_Ratio']!  
    →=0)&(transformed_data['Age']==28))]['Superplasticizer_Ratio']  
strength = transformed_data[((transformed_data['Superplasticizer_Ratio']!  
    →=0)&(transformed_data['Age']==28))]['Compressive_Strength']  
plt.scatter(superplasticizer,strength)
```

```
[68]: <matplotlib.collections.PathCollection at 0x7ff5606ed580>
```



This is better, but we see a large spread in the data. Let's remove any outliers first, before training our model.

```
[69]: superplasticizer.describe()
```

```
[69]: count    317.000000  
     mean      0.004031  
     std      0.001713  
     min      0.000746  
     25%      0.002988  
     50%      0.003910
```

```
75%          0.004776
max          0.013149
Name: Superplasticizer_Ratio, dtype: float64
```

```
[70]: mean = 0.004146
three_sigma = 3*0.001875
upper = mean + three_sigma
lower = mean - three_sigma


print(f"The lower bound is:\t{lower}\n\nThe upper bound is:\t{upper}")
```

```
The lower bound is:      -0.0014789999999999994
The upper bound is:      0.009771
```

Since there are no negative ratios, we only need to remove data points where the superplasticizer ratio is greater than 0.009771.

```
[71]: superplasticizer = transformed_data[
    ((transformed_data['Superplasticizer_Ratio'] != 0) &
    (transformed_data['Age'] == 28) &
    (transformed_data['Superplasticizer_Ratio'] < upper))]['Superplasticizer_Ratio']
strength = transformed_data[
    ((transformed_data['Superplasticizer_Ratio'] != 0) &
    (transformed_data['Age'] == 28) &
    (transformed_data['Superplasticizer_Ratio'] < upper))]['Compressive Strength']
plt.scatter(superplasticizer, strength)
```

```
[71]: <matplotlib.collections.PathCollection at 0x7ff5607c6730>
```



1.11.2 Train the Linear Model

```
[72]: # We will train and test our model only on the data above, that does not
      ↪ contain outliers
      # Reshape the data so it complies with the linear model requirements
      X = np.array(superplasticizer).reshape(315,1)
      y = np.array(strength).reshape(315,1)
```

```
[73]: # Perform a train-test split
      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
      ↪ 2,random_state=42)

      # Train the linear model
      lm = LinearRegression()
      lm.fit(X_train,y_train)
```

```
[73]: LinearRegression()
```

1.11.3 Test the Linear Model

```
[74]: y_pred = lm.predict(X_test)
```

1.11.4 Linear Equation

```
[75]: # print the intercept
      print(lm.intercept_)
```

```
[36.83431501]
```


```
[76]: coeff = pd.DataFrame(lm.coef_,columns=['Coefficient'])
      coeff
```

```
[76]:    Coefficient
      0    273.960172
```

1.11.5 Model Evaluation

```
[77]: # Plot the linear model predictions as a line superimposed on a scatter plot of
      ↪ the testing data
      plt.scatter(X_test,y_test)
      plt.plot(X_test,y_pred,'r')
```

```
[77]: [<matplotlib.lines.Line2D at 0x7ff5608cbe50>]
```



```
[78]: # Evaluation Metrics
MAE_super_28 = mean_absolute_error(y_test, y_pred)
MSE_super_28 = mean_squared_error(y_test, y_pred)
RMSE_super_28 = np.sqrt(mean_squared_error(y_test, y_pred))

super_stats_28 = [MAE_super_28, MSE_super_28, RMSE_super_28] # storing for model_
↳ comparison at the end of this notebook

# Print the metrics
print(f"EVALUATION METRICS, LINEAR MODEL FOR SUPERPLASTICIZER RATIO VS.
↳ COMPRESSIVE STRENGTH AT 28 DAYS")
print('-----')
print(f"Mean Absolute Error (MAE):\t\t{MAE_super_28}\nMean Squared Error:
↳ \t\t\t{MSE_super_28}\nRoot Mean Squared Error (RMSE):\t\t{RMSE_super_28}")
print('-----\n\n')
```

EVALUATION METRICS, LINEAR MODEL FOR SUPERPLASTICIZER RATIO VS. COMPRESSIVE STRENGTH AT 28 DAYS

```
-----
Mean Absolute Error (MAE):                11.954354128234902
Mean Squared Error:                      243.9703785891613
Root Mean Squared Error (RMSE):          15.619551164779393
-----
```


1.12 Model Comparisons Analysis

Neither superplasticizer linear model appeared to represent the data well from a visual perspective. The cement, cementitious ratio, and fly ash ratio linear models, however, did. We can display all of the evaluation metrics below and compare them to the artificial neural network's (ANN) performance.

```
[103]: ANN_metrics = [5.083552,6.466492**2,6.466492]

metrics = [cement_stats, cementitious_stats, fly_stats, super_stats,
           ↪ANN_metrics]
metrics_28 = [cement_28_stats, cementitious_28_stats, fly_28_stats,
             ↪super_stats_28, ANN_metrics]

metrics_df = pd.DataFrame(data=metrics, index=['Cement (Ignoring Cure
           ↪Time)', 'Cementitious_Ratio (Ignoring Cure Time)', 'Fly_Ash_Ratio (Ignoring
           ↪Cure Time)', 'Superplasticizer_Ratio (Ignoring Cure Time)', 'ANN (Function of
           ↪Time)'], columns=['MAE', 'MSE', 'RMSE'])
metrics_28_df = pd.DataFrame(data=metrics_28, index=['Cement (Cure Time = 28
           ↪Days)', 'Cementitious_Ratio (Cure Time = 28 Days)', 'Fly_Ash_Ratio (Cure Time
           ↪= 28 Days)', 'Superplasticizer_Ratio (Cure Time = 28 Days)', 'ANN (Function of
           ↪Time)'], columns=['MAE', 'MSE', 'RMSE'])
```

```
[104]: metrics_df
```

```
[104]:
```

	MAE	MSE	RMSE
Cement (Ignoring Cure Time)	11.555613	192.784799	13.884697
Cementitious_Ratio (Ignoring Cure Time)	12.834673	253.058148	15.907801
Fly_Ash_Ratio (Ignoring Cure Time)	12.121987	212.894325	14.590899
Superplasticizer_Ratio (Ignoring Cure Time)	12.189767	225.909298	15.030279
ANN (Function of Time)	5.083552	41.815519	6.466492

```
[105]: metrics_28_df
```

```
[105]:
```

	MAE	MSE	RMSE
Cement (Cure Time = 28 Days)	9.134082	140.117503	11.837124
Cementitious_Ratio (Cure Time = 28 Days)	11.519580	197.957201	14.069726
Fly_Ash_Ratio (Cure Time = 28 Days)	7.692695	91.317960	9.556043
Superplasticizer_Ratio (Cure Time = 28 Days)	11.954354	243.970379	15.619551
ANN (Function of Time)	5.083552	41.815519	6.466492

1.13 Conclusions & Recommendations

By comparing the evaluation metrics for all models, we conclude that the ANN model performed significantly better than all of the linear models. It outperformed the best linear model's RMSE (for Fly_Ash_Ratio at 28 Days) by over 30%! An important note is that the linear models were not scaled, and the ANN model was. We kept the linear models biased in order to maintain coefficient interpretability, whereas that was not relevant to the ANN model.

What is surprising is that the ANN model still outperformed the linear models, even when controlling for cure time at 28 days. Perhaps the most startling insight is that the fly ash ratio was even more accurate at predicting concrete compressive strength than the cement quantity, to the point that it had the lowest errors of all of the linear models. We therefore recommend that engineers give very conservative fly ash ratio specifications when allowing substitutions for Portland cement.